

# Email Search AI: A Generative Search System for Enterprise Email Intelligence

## Abstract

Organizations rely heavily on email communication to make strategic decisions, approve budgets, and define execution timelines. However, as email volume grows, extracting accurate, decision-oriented insights from large corpora becomes increasingly difficult. This project presents **Email Search AI**, a robust **generative search system** built using modern embedding models, vector databases, re-ranking techniques, and large language models (LLMs). The system enables users to query large collections of enterprise emails using natural language and receive concise, fact-based answers grounded in retrieved evidence.

The system follows a **three-layer architecture**: Embedding Layer, Search Layer, and Generation Layer. Extensive experimentation was conducted on chunking strategies, embedding models, re-ranking approaches, and prompt design to optimize retrieval accuracy and answer quality.

## 1. Introduction

### 1.1 Problem Statement

Enterprise emails often contain:

- Strategic decisions
- Budget approvals
- Project timelines
- Action items and responsibilities

Traditional keyword-based email search systems fail to:

- Capture semantic meaning
- Aggregate information across long threads
- Distinguish finalized decisions from discussions

This results in lost institutional knowledge and inefficient decision validation.

## 1.2 Project Objective

The goal of this project is to build a **semantic, generative search system** that:

- Retrieves relevant email content using semantic similarity
- Re-ranks results for precision
- Generates concise, factual answers using an LLM
- Grounds responses in verifiable email evidence

## 2. Dataset Description

### 2.1 Data Sources

The project uses a publicly available email dataset consisting of:

#### 1. *email\_thread\_details.csv*

Contains:

- `thread_id`
- `subject`
- `timestamp`
- `from`
- `to`
- `body`

This file represents raw organizational email communication.

#### 2. *email\_thread\_summaries.csv*

Contains:

- `thread_id`
- `summary`

These summaries provide high-level context for entire email threads and are used to enhance retrieval and generation.

## 2.2 Data Cleaning


Email bodies often contain redundant reply chains and metadata. The following preprocessing steps were applied:

- Removal of reply headers (e.g., “On X wrote:”)
- Removal of inline metadata (From, To, Subject)
- Normalization of whitespace

This significantly improves embedding quality by reducing noise.

## 3. System Architecture Overview

The system is built using a modular **three-layer architecture**, enabling experimentation and scalability.



## 4. Embedding Layer

### 4.1 Purpose

The embedding layer converts unstructured email text into dense vector representations that capture semantic meaning.

### 4.2 Chunking Strategy

#### *Strategy Used*

A **word-based sliding window chunking strategy** was implemented:

- Chunk size: **300 words**
- Overlap: **50 words**

#### *Motivation*

- Emails often contain long, multi-paragraph discussions
- Decisions may span multiple paragraphs
- Overlapping chunks preserve contextual continuity

#### *Alternatives Considered*

Strategy	Limitations
Sentence-based chunking	Breaks long decisions
Fixed paragraph chunking	Inconsistent size
No overlap	Loss of boundary context

The selected strategy provided the best balance between **context preservation and retrieval accuracy**.

## 4.3 Embedding Model Selection

### Model Used:

SentenceTransformer("all-MiniLM-L6-v2")

### *Justification:*

- Optimized for semantic search
- Efficient inference for large datasets
- Strong performance on retrieval benchmarks
- Lower computational cost than larger transformer models

Both **email chunks** and **thread summaries** were embedded using the same model to ensure embedding space consistency.

## 4.4 Metadata Design

Each embedded document stores structured metadata:

- Thread ID
- Subject
- Timestamp
- Sender and recipient
- Document type (`email` or `thread_summary`)

This metadata is later used in:

- Re-ranking
- Prompt construction
- Evidence presentation

## 5. Search Layer

### 5.1 Vector Database

**ChromaDB** was used as the vector store with persistent storage enabled.

#### *Advantages:*

- Fast similarity search
- Metadata filtering
- Simple integration with Python
- Persistent local storage for reproducibility

### 5.2 Query Processing

Each user query is:

1. Embedded using the same MiniLM model
2. Searched against the vector database
3. Retrieved with top-K semantic matches

Initial retrieval prioritizes **recall** over precision.

### 5.3 Cache Implementation

A query-level cache was implemented using MD5 hashing.

#### *Benefits:*

- Prevents redundant embedding and retrieval
- Improves response time for repeated queries
- Reduces computational cost

## 5.4 Re-ranking Mechanism

### *Model Used:*

cross-encoder/ms-marco-MiniLM-L-6-v2

### *Why Re-ranking?*

Vector similarity alone may retrieve:

- Semantically related but irrelevant content
- Contextually incomplete chunks

The cross-encoder jointly encodes (query, document) pairs to:

- Improve ranking precision
- Identify decision-critical content

### *Final Output:*

Top **3 highly relevant chunks** per query.

## 6. Generation Layer

### 6.1 Prompt Engineering

The prompt is structured into distinct sections:

- Thread summaries
- Email evidence with timestamps and participants
- Explicit question
- Instruction for concise, factual output

This design:

- Grounds the LLM in retrieved evidence
- Reduces hallucinations

- Encourages decision-focused answers

## 6.2 Language Model

**Model Used:** gpt-4o-mini

**Configuration:**

- Temperature: 0
- Deterministic responses
- Optimized for factual summarization

## 6.3 Answer Characteristics

- Concise
- Fact-based
- Evidence-driven
- Mentions timelines, approvals, and decisions explicitly

# 7. Evaluation

## 7.1 Self-Designed Queries

1. *What decisions were made about Q3 marketing strategy?*
2. *Was budget approval discussed for Project Atlas?*
3. *What timelines were agreed upon for product launch?*

## 7.2 Observations

- Correct retrieval of decision-oriented emails
- Accurate aggregation across multiple threads

- Clear differentiation between discussions and finalized decisions

## 8. Challenges & Limitations

### Challenges

- Email noise from long reply chains
- Implicit decisions without explicit language
- Temporal ambiguity in discussions

### Limitations

- No date-range filtering
- No named entity normalization
- Relies on quality of input summaries

## 9. Future Work

- Time-aware search (before/after date filters)
- Entity recognition for projects and stakeholders
- UI for evidence highlighting
- Fine-tuned domain-specific embedding models

## 10. Conclusion

Email Search AI demonstrates how **modern LLM-powered search systems** can unlock institutional knowledge hidden in enterprise emails. By combining semantic embeddings, vector search, re-ranking, and structured generation, the system delivers accurate, explainable, and decision-focused answers. This architecture is extensible and well-suited for real-world organizational deployment.